

```
-- IMPORTANT: BEFORE CREATING ANY TABLE, MAKE SURE YOU RUN THIS COMMAND
ADD JAR /opt/cloudera/parcels/CDH/lib/hive/lib/hive-hcatalog-core-1.1.0-cdh5.11.2.jar;
```

```
-- Drop table if exists
DROP TABLE nyc_taxi_data;
```

```
-- create an external table
```

```
CREATE EXTERNAL TABLE IF NOT EXISTS nyc_taxi_data(
  `vendorid` int,
  `tpep_pickup_datetime` timestamp,
  `tpep_dropoff_datetime` timestamp,
  `passenger_count` int,
  `trip_distance` double,
  `ratecodeid` int,
  `store_and_fwd_flag` string,
  `pulocationid` int,
  `dolocationid` int,
  `payment_type` int,
  `fare_amount` double,
  `extra` double,
  `mta_tax` double,
  `tip_amount` double,
  `tolls_amount` double,
  `improvement_surcharge` double,
  `total_amount` double)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
STORED AS TEXTFILE
LOCATION '/common_folder/nyc_taxi_data'
TBLPROPERTIES ("skip.header.line.count"="1");
```

```
-- Check for the created tables
```

```
select * FROM nyc_taxi_data limit 10;
describe nyc_taxi_data
```

```
-- BASIC DATA QUALITY CHECKS
```

```
-----
-----
```

```
--          1. How many records has each TPEP provider provided?
--          Write a query that summarises the number of records of each
provider.
```

```
-----
-----
```

```
SELECT vendorid,count(*) AS count
FROM nyc_taxi_data
GROUP BY vendorid;
```

```
-----
--          vendorid      count
--          1              527386
--          2              647183
-----
```

```
-----
-----
```

```

--          2. The data provided is for months November and December
only.
--          Check whether the data is consistent, and if not, identify
the data quality issues.
--          Mention all data quality issues in comments.
-----
-----

-- DATETIME ISSUES

-- Analyzing pickup time

SELECT  vendorid, YEAR(tpep_pickup_datetime) AS yr,
MONTH(tpep_pickup_datetime) AS mnth, count(*) AS total
FROM nyc_taxi_data
GROUP BY vendorID, YEAR(tpep_pickup_datetime), MONTH(tpep_pickup_datetime)
ORDER BY vendorID, yr, mnth;

-- Analyzing drop time

SELECT  vendorid, YEAR(tpep_dropoff_datetime) AS yr,
MONTH(tpep_dropoff_datetime) AS mnth, count(*) AS total
FROM nyc_taxi_data
GROUP BY vendorID, YEAR(tpep_dropoff_datetime),
MONTH(tpep_dropoff_datetime)
ORDER BY vendorID, yr, mnth;

-----
-----

--          There are data with years 2003, 2008, 2009, 2018, 2019
--          Data with 2018 with month 1 can be acceptable as the trip has
started in 31 DEC 2017 and ended up in 1 Jan 2018
--          But Data with other years with random months are erroneous and
vendor 2 is providing almost all of them.

-----
-----

--          3. You might have encountered unusual or erroneous rows in
the dataset.
--          Can you conclude which vendor is doing a bad job in
providing the records?
-----
-----

--          We have only 2 Vendor IDs. To review the job done by the vendors,
we can use the rate code IDs
--          Analyzing the Rate code IDs

SELECT  vendorid, ratecodeid, count(*) as count
FROM nyc_taxi_data
GROUP BY vendorid, ratecodeid
ORDER BY vendorid, ratecodeid;

-----
-----

--          vendorid          ratecodeid          _c2
--          1                  1                  513991

```

--	1	2	10544
--	1	3	1186
--	1	4	230
--	1	5	1425
--	1	6	2
--	1	99	8
--	2	1	628287
--	2	2	14794
--	2	3	1376
--	2	4	356
--	2	5	2368
--	2	6	1
--	2	99	1

-- Both the vendors are not performing a good job as the rating is 1 for a huge share of trips
 -- Also, both the vendors are seeding invalid ratecodeIDs,
 -- in which vendor ID 1 (Creative Mobile Technologies) is top with 8 records.

-- UNUSUAL OR ERRONEOUS PASSENGER COUNT

```
SELECT vendorid, passenger_count, count(*) AS count
FROM nyc_taxi_data
GROUP BY vendorid, passenger_count
ORDER BY vendorID, passenger_count;
```


-- Vendor 1 and 2 are seeding unusual passenger_count i.e equal to 0
 -- The number for Vendor 1 is higher compared to Vendor 2

-- UNUSUAL OR ERRONEOUS FARE AMOUNT

```
SELECT vendorid, count(*) AS count_error, sum(fare_amount) AS sum_error
FROM nyc_taxi_data
WHERE fare_amount<0
GROUP BY vendorid;
```


--	vendorid	count_error	sum_error	
--	2	558		-4917.38

-- fare_amount was unusually charged by vendor 2, i.e 558 times with total of -\$4917(negative)

-- UNUSUAL OR ERRONEOUS EXTRA CHARGE

```
SELECT vendorid, count(*) AS count_error, sum(extra) AS sum_error
FROM nyc_taxi_data
WHERE extra NOT IN (0,0.5,1)
GROUP BY vendorid;
```


-- Extra charges are also unusual somewhere, where Both are charging
some unusual Extra as 1823 and 3033
-- times respectively, where vendor 2 (VeriFone Inc) has done this
more times with approx. 4000 USD

-- UNUSUAL OR ERRONEOUS MTA TAX

```
SELECT vendorid, count(*) AS count_error, sum(mta_tax) AS sum_error
FROM nyc_taxi_data
WHERE mta_tax NOT IN (0,0.5)
GROUP BY vendorid;
```


-- mta_tax was also unusually charged, mostly from vendor 2, i.e 547
times with total of -\$263(negative)

-- UNUSUAL OR ERRONEOUS TIP AMOUNT

```
SELECT vendorid, count(*) AS count_error, sum(tip_amount) AS sum_error
FROM nyc_taxi_data
WHERE tip_amount<0
GROUP BY vendorid;
```


-- vendorid count_error sum_error
-- 2 4 -3.5

-- tip_amount was unusually charged in a few cases, all from vendor 2, i.e
4 times with total of -\$3.5 (negative)

-- UNUSUAL OR ERRONEOUS TIP AMOUNT IMPROVEMENT_SURCHARGE

```
SELECT vendorid, count(*) AS count_error, sum(improvement_surcharge) AS
sum_error
FROM nyc_taxi_data
```

```
WHERE improvement_surcharge NOT IN (0,0.3)
GROUP BY vendorid;
```

```
-----
-----
--      improvement_surcharge was also unusually charged, all from vendor
2, i.e 562 times with total of -$163.4 (negative)
-----
-----
```

```
-- UNUSUAL OR ERRONEOUS PICKUP AND DROP TIME
```

```
SELECT vendorid ,count(*) AS count_error
FROM nyc_taxi_data
WHERE
UNIX_TIMESTAMP(tpep_pickup_datetime)>UNIX_TIMESTAMP(tpep_dropoff_datetime)
GROUP BY vendorid;
```

```
-----
-----
--      vendorid      count_error
--      1              73
```

```
-- Vendor 1 is seeding unusual pickup time i.e greater than drop time
-----
-----
```

```
-- UNUSUAL OR ERRONEOUS TRIP DURATION
```

```
SELECT vendorid ,count(*) as count_error
FROM nyc_taxi_data
WHERE trip_distance <= 0
GROUP BY vendorid;
```

```
-----
-----
--      vendorid      count_error
--      1              3185
--      2              4217
-----
-----
```

```
----- Final Conclusion on unusual or erroneous rows in the
dataset -----
```

```
-- There are unusual or erroneous rows in dataset with respect to:
```

```
-- RATE CODES
-- PASSENGER COUNT
-- FARE AMOUNT
-- TRIP DURATION
-- PASSENGER COUNT
-- PICKUP & DROP TIME
-- MTA TAX
-- EXTRA CHARGE
```

```

-- TIP AMOUNT
-- FARE AMOUNT
-- IMPROVEMENT SURCHARGE

--      Detail about these unusual or erroneous rows has been discussed
above with the queries

--      Note: Both 1 and 2, are correct answers as per the accountability
of the erroneous points

-----
-----

-- CREATING A CLEAN, ORC PARTITIONED TABLE FOR ANALYSIS

-- CREATING A PARTITIONED TABLE WITH FORMATTED DATA AND IN ORC COMPRESSION
-- IMPORTANT: BEFORE PARTITIONING ANY TABLE, MAKE SURE YOU RUN THESE
COMMANDS

SET hive.exec.max.dynamic.partitions=100000;
SET hive.exec.max.dynamic.partitions.pernode=100000;

-- Drop table if already exists
DROP TABLE IF EXISTS nyc_taxi_data_partitioned_orc;

-- Creating table with required datatypes(columns), partition settings and
compressd format configuration

CREATE EXTERNAL TABLE IF NOT EXISTS nyc_taxi_data_partitioned_orc(
`vendorid` int,
`tpep_pickup_datetime` timestamp,
`tpep_dropoff_datetime` timestamp,
`passenger_count` int,
`trip_distance` double,
`ratecodeid` int,
`store_and_fwd_flag` string,
`pulocationid` int,
`dolocationid` int,
`payment_type` int,
`fare_amount` double,
`extra` double,
`mta_tax` double,
`tip_amount` double,
`tolls_amount` double,
`improvement_surcharge` double,
`total_amount` double
)
PARTITIONED BY (yr int, mnth int)
STORED AS ORC
LOCATION '/user/hive/warehouse/nyc_taxi_data_partitioned_orc'
TBLPROPERTIES ("orc.compress"="SNAPPY");

SELECT * FROM nyc_taxi_data_partitioned_orc;

-- Setting to allow partition limits and insert permission into partition
SET hive.exec.dynamic.partition=true;
SET hive.exec.dynamic.partition.mode=nonstrict;

-- Inserting data to the partitioned table
INSERT OVERWRITE TABLE nyc_taxi_data_partitioned_orc PARTITION(yr, mnth)

```

```

SELECT `vendorid` int,
`tpep_pickup_datetime` timestamp,
`tpep_dropoff_datetime` timestamp,
`passenger_count` int,
`trip_distance` double,
`ratecodeid` int,
`store_and_fwd_flag` string,
`pulocationid` int,
`dolocationid` int,
`payment_type` int,
`fare_amount` double,
`extra` double,
`mta_tax` double,
`tip_amount` double,
`tolls_amount` double,
`improvement_surcharge` double,
`total_amount` double,
YEAR(`tpep_pickup_datetime`) AS yr,
MONTH(`tpep_pickup_datetime`) AS mnth
FROM nyc_taxi_data
WHERE YEAR(`tpep_pickup_datetime`) = 2017 AND MONTH(`tpep_pickup_datetime`)
in (11,12)
AND `passenger_count` NOT IN (0, 192)
AND `ratecodeid` IN (1,2,3,4,5,6);

```

```

-----
--      Note: We have removed the erroneous data pertaining to columns:
tpep_pickup_datetime, passenger_count and ratecodeid

```

```

-----
--      Checking for data in table

```

```

SELECT * FROM nyc_taxi_data_partitioned_orc LIMIT 10;

```

```

--      Checking for total records available

```

```

SELECT vendorid, COUNT(*) AS count
FROM nyc_taxi_data_partitioned_orc
GROUP BY vendorid
ORDER BY vendorid;

```

```

-----
--      vendorid      count (previous)
--      1              520571 (527386)
--      2              647157 (647183)
-----

```

```

-- Note: Here we have used partitions, buckets and ORC AS well
-- We don't expect the same from student, even if they create "Partitions"
and "ORC" individually, we'll provide them full marks

```

```

-----
-- ANALYSIS-I

```

```
-----  
-----  
-- 1. Compare the average fare for November and December.  
-----  
-----
```

```
SELECT mnth, round(avg(fare_amount),2) AS Average_fare_amount  
FROM nyc_taxi_data_partitioned_orc  
WHERE fare_amount>0  
GROUP BY mnth;
```

```
-----  
-----  
--      mnth      average_fare_amount  
--      11         13.12  
--      12         12.91
```

```
--      Average Fare amount in November is slightly higher than the same in  
December.
```

```
-----  
-----  
--      2. Explore the number of passengers per trip - how many trips are  
made by each level  
--      of Passenger_count? Do most people travel solo or with other  
people?  
-----  
-----
```

```
SELECT passenger_count, count(*) AS Total  
FROM nyc_taxi_data_partitioned_orc  
WHERE isnotnull(passenger_count) and passenger_count>0  
GROUP BY passenger_count  
ORDER BY passenger_count;
```

```
-----  
-----  
--      passenger_count      total  
--      1                     827485  
--      2                     176871  
--      3                     50693  
--      4                     24951  
--      5                     54567  
--      6                     33145  
--      7                      12  
--      8                      3  
--      9                      1
```

```
--      Most of the people travel solo, as we can see the count is highest  
(827485) when passenger_count is 1
```

```
-----  
-----  
-- 3. Which is the most preferred mode of payment?  
-----  
-----
```

```
SELECT payment_type, count(*) AS Total  
FROM nyc_taxi_data_partitioned_orc
```



```
GROUP BY payment_type;
```

```
-----  
-----
```

```
--      payment_mode          total  
--      Credit card (1)      785728  
--      Cash (2)              374179  
--      No charge (3)        6187  
--      Dispute (4)           1634  
  
--      Credit Card is most preffered mode of payment
```

```
-----  
-----
```

```
-- 4. What is the average tip paid? Compare the average tip with the 25th,  
50th and 75th percentiles  
-- and comment whether the average tip is a representative statistic (of  
the central tendency)  
-- of tip amount paid.
```

```
-----  
-----
```

```
SELECT AVG(tip_amount) AS average_tip  
FROM nyc_taxi_data_partitioned_orc  
WHERE tip_amount >= 0  
AND fare_amount > 0;
```

```
-- Average tip_amount paid  
--      1.8545392085777652
```

```
SELECT PERCENTILE_APPROX(`tip_amount`, 0.25) AS 25_per,  
PERCENTILE_APPROX(`tip_amount`, 0.50) AS 50_Per,  
PERCENTILE_APPROX(`tip_amount`, 0.75) AS 75_per,  
AVG(`tip_amount`) AS Average  
FROM nyc_taxi_data_partitioned_orc  
WHERE `tip_amount` >= 0  
AND `fare_amount` > 0;
```

```
-----  
-----
```

```
--      25th, 50th and 75th percentiles of tip amount  
--      [0.0,1.36,2.45]
```

```
-----  
-----
```

```
-- Average tip_amount paid      : 1.85  
-- 50th percentile of tip amount : 1.36  
-- The average tip is not a representative statistic (of the central  
tendency) of tip amount paid .
```

```
-----  
-----
```

```
-- 5. Explore the Extra (charge) variable - what is the fraction of total  
trips where an extra charge is levied?
```

```
-----  
-----
```

```
SELECT SUM(IF( extra > 0, 1 , 0 ))/ COUNT(*) * 100 AS  
Fraction_When_Levied_Extra  
FROM nyc_taxi_data_partitioned_orc
```

```

-----
-----

--      46.21% of trips happen when extra charge is levied, means people
like to travel when no extra charge is levied.

-----
-----

-- ANALYSIS-II

-----
-----

-- 1.What is the correlation between the number of passengers and tip paid?
Do multiple travellers
-- pay more compared to solo travellers?
-----
-----

-- Finding correlation

SELECT CORR(tip_amount, passenger_count)
from nyc_taxi_data_partitioned_orc
WHERE tip_amount>=0 AND passenger_count>0

----- OUTPUT -----
--      -0.0047482397788550035
-----

-- Verifying correlation by Correlation Coefficient( $r = \text{Cov}(x,y) / S_x * S_y$ )

SELECT covar_pop(tip_amount,
passenger_count)/(stddev_pop(tip_amount)*stddev_pop(passenger_count))
from nyc_taxi_data_partitioned_orc
WHERE tip_amount>=0 AND passenger_count>0

----- OUTPUT -----
--      -0.00474823977885501
-----

-- CONCLUSION

-- Correlation between the number of passengers and tip paid: -
0.0047482397788550035
-- It indicates Weak Negative Correlation.
-- It means as number of passengers increases, the tip amount decreases
very slightly.
-- Based on correlation value, solo travellers pay more compared to
multiple travellers.

-----
-----

--      Q2. Segregate the data into five segments of "tip paid": [0-5),
[5-10), [10-15) , [15-20) and >=20.
--      Calculate the percentage share of each bucket (i.e. the fraction of
trips falling in each bucket).
-----
-----

-- Below query returns the fraction of trips falling in each segment.

```

```

SELECT (SUM(IF(tip_amount >=0 AND tip_amount < 5, 1,0))/COUNT(*))*100 AS
`[0-5)` ,
      (SUM(IF(tip_amount >=5 AND tip_amount < 10, 1,0))/COUNT(*))*100 AS
`[5-10)` ,
      (SUM(IF(tip_amount >=10 AND tip_amount < 15, 1,0))/COUNT(*))*100 AS
`[10-15)` ,
      (SUM(IF(tip_amount >=15 AND tip_amount < 20, 1,0))/COUNT(*))*100 AS
`[15-20)` ,
      (SUM(IF(tip_amount >=20, 1,0))/COUNT(*))*100 AS `>=20`
FROM nyc_taxi_data_partitioned_orc
WHERE tip_amount >= 0
AND fare_amount > 0;

```

-- CONCLUSION

```

-- Fraction of Trips Falling in Bucket [0-5)      - 92.11106369047926
-- Fraction of Trips Falling in Bucket [5-10)     - 5.654481849610322
-- Fraction of Trips Falling in Bucket [10-15)    - 1.8930259105189817
-- Fraction of Trips Falling in Bucket [15-20)    - 0.23687462613402444
-- Fraction of Trips Falling in Bucket >=20       - 0.10455392325742033

```

```

-----
--      Q3. Which month has a greater average speed - November or
December? Note that the
--      variable speed will have to be derived from other metrics.
-----

```

-- Below query returns average speed for November & December 2017.

```

SELECT mnth,
AVG(trip_distance/((UNIX_TIMESTAMP(tpep_dropoff_datetime) -
UNIX_TIMESTAMP(tpep_pickup_datetime))/3600)) as Avg_Speed_MPH
FROM nyc_taxi_data_partitioned_orc
WHERE trip_distance >= 0
GROUP BY mnth;

```

-- CONCLUSION

```

-----
-- November Month has average speed as 10.97802043563046 Miles Per Hour
-- December Month has average speed as 11.073593998600314 Miles Per Hour
-- Based on average speed values, December Month has a greater average
speed

```

```

-----
--      Q4. Analyse the average speed of the most happening days of the
year i.e. 31st December
--      (New year's eve) and 25th December (Christmas Eve) and compare it
with the overall average.
-----

```

-- Below query returns overall average speed for both November & December 2017.

```

SELECT AVG(trip_distance/((UNIX_TIMESTAMP(tpep_dropoff_datetime) -
UNIX_TIMESTAMP(tpep_pickup_datetime))/3600)) as Avg_Speed_MPH

```

```
FROM nyc_taxi_data_partitioned_orc
WHERE trip_distance >= 0
AND YEAR(tpep_dropoff_datetime) IN (2017, 2018);
```

-- overall average speed of the trips: 11.026369911646409 Miles Per Hour

-- Below query returns average speed on 31st December 2017 & 25th December 2017.

```
SELECT FROM_UNIXTIME(UNIX_TIMESTAMP(tpep_pickup_datetime), 'dd-MMM-yyyy')
as `Happening_date`,
AVG(trip_distance/((UNIX_TIMESTAMP(tpep_dropoff_datetime) -
UNIX_TIMESTAMP(tpep_pickup_datetime))/3600)) as Avg_Speed_MPH
FROM nyc_taxi_data_partitioned_orc
WHERE trip_distance >= 0
AND mnth = 12
AND DAY(tpep_pickup_datetime) IN (25,31)
AND YEAR(tpep_dropoff_datetime) IN (2017, 2018)
GROUP BY FROM_UNIXTIME(UNIX_TIMESTAMP(tpep_pickup_datetime), 'dd-MMM-
yyyy');
```


-- CONCLUSION

-- On 25-Dec-2017, the average speed was 15.24030794591516 Miles Per Hour
-- On 31-Dec-2017, the average speed was 13.202755584924587 Miles Per Hour
-- Overall average speed was 11.026369911646409
-- Based on average speed values analysis, the average speed was vewry high on 25-Dec-2017.

